

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import warnings
import shap
from sklearn.preprocessing import OneHotEncoder
from category_encoders import TargetEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from xgboost import XGBRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split, RandomizedSearchCV, cross_val_score, KFold
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv("/kaggle/input/airlines-flights-data/airlines_flights_data.csv")
df.head()
```

```
Out[2]:
```

	index	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
0	0	SpiceJet	SG-8709	Delhi	Evening	zero	Night	Mumbai	Economy	2.17	1	5953
1	1	SpiceJet	SG-8157	Delhi	Early_Morning	zero	Morning	Mumbai	Economy	2.33	1	5953
2	2	AirAsia	I5-764	Delhi	Early_Morning	zero	Early_Morning	Mumbai	Economy	2.17	1	5956
3	3	Vistara	UK-995	Delhi	Morning	zero	Afternoon	Mumbai	Economy	2.25	1	5955
4	4	Vistara	UK-963	Delhi	Morning	zero	Morning	Mumbai	Economy	2.33	1	5955

```
In [3]: df.drop(columns=['index'], inplace=True)
```

## EXPLORATORY DATA ANALYSIS

```
In [4]: df.head()
```

```
Out[4]:
```

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
0	SpiceJet	SG-8709	Delhi	Evening	zero	Night	Mumbai	Economy	2.17	1	5953
1	SpiceJet	SG-8157	Delhi	Early_Morning	zero	Morning	Mumbai	Economy	2.33	1	5953
2	AirAsia	I5-764	Delhi	Early_Morning	zero	Early_Morning	Mumbai	Economy	2.17	1	5956
3	Vistara	UK-995	Delhi	Morning	zero	Afternoon	Mumbai	Economy	2.25	1	5955
4	Vistara	UK-963	Delhi	Morning	zero	Morning	Mumbai	Economy	2.33	1	5955

```
In [5]: df.describe()
```

```
Out[5]:
```

	duration	days_left	price
count	300153.000000	300153.000000	300153.000000
mean	12.221021	26.004751	20889.660523
std	7.191997	13.561004	22697.767366
min	0.830000	1.000000	1105.000000
25%	6.830000	15.000000	4783.000000
50%	11.250000	26.000000	7425.000000
75%	16.170000	38.000000	42521.000000
max	49.830000	49.000000	123071.000000

```
In [6]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300153 entries, 0 to 300152
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   airline                300153 non-null object  
1   flight                 300153 non-null object  
2   source_city            300153 non-null object  
3   departure_time         300153 non-null object  
4   stops                  300153 non-null object  
5   arrival_time           300153 non-null object  
6   destination_city       300153 non-null object  
7   class                  300153 non-null object  
8   duration                300153 non-null float64  
9   days_left              300153 non-null int64   
10  price                  300153 non-null int64   
dtypes: float64(1), int64(2), object(8)
memory usage: 25.2+ MB

```

```
In [7]: df[df['duration'] == df['duration'].max()]
```

```
Out[7]:
```

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
<b>193889</b>	Air_India	AI-672	Chennai	Evening	two_or_more	Evening	Bangalore	Economy	49.83	2	23891
<b>194359</b>	Air_India	AI-672	Chennai	Evening	one	Evening	Bangalore	Economy	49.83	9	17538

```
In [8]: df[df['duration'] == df['duration'].min()]
```

Out[8]:

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
115869	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	6	3498
115943	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	7	3498
116010	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	8	3498
116081	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	9	3498
116163	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	10	3498
116236	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	11	3498
116322	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	12	3498
116411	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	13	3498
116496	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	14	3498
116656	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	16	1924
116835	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	18	1924
116924	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	19	1924
117019	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	20	1924
117101	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	21	1924
117190	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	22	1924
117274	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	23	1924
117366	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	24	1924
117461	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	25	1924
117547	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	26	1924
117643	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	27	1924
117728	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	28	1924
117817	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	29	1924

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
117900	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	30	1924
117995	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	31	1604
118086	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	32	1604
118173	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	33	1604
118269	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	34	1604
118355	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	35	1604
118445	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	36	1604
118528	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	37	1604
118622	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	38	1604
118712	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	39	1604
118799	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	40	1604
118896	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	41	1604
118982	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	42	1604
119072	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	43	1604
119155	Indigo	6E-357	Bangalore	Night	zero	Night	Chennai	Economy	0.83	44	1604
197354	Indigo	6E-987	Chennai	Early_Morning	zero	Early_Morning	Bangalore	Economy	0.83	45	1443
197355	Indigo	6E-477	Chennai	Early_Morning	zero	Morning	Bangalore	Economy	0.83	45	1443
197356	Indigo	6E-6137	Chennai	Morning	zero	Morning	Bangalore	Economy	0.83	45	1443
197445	Indigo	6E-987	Chennai	Early_Morning	zero	Early_Morning	Bangalore	Economy	0.83	46	1443
197446	Indigo	6E-477	Chennai	Early_Morning	zero	Morning	Bangalore	Economy	0.83	46	1443
197447	Indigo	6E-6137	Chennai	Morning	zero	Morning	Bangalore	Economy	0.83	46	1443
197537	Indigo	6E-987	Chennai	Early_Morning	zero	Early_Morning	Bangalore	Economy	0.83	47	1443

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
<b>197538</b>	Indigo	6E-477	Chennai	Early_Morning	zero	Morning	Bangalore	Economy	0.83	47	1443
<b>197539</b>	Indigo	6E-6137	Chennai	Morning	zero	Morning	Bangalore	Economy	0.83	47	1443
<b>197626</b>	Indigo	6E-987	Chennai	Early_Morning	zero	Early_Morning	Bangalore	Economy	0.83	48	1443
<b>197627</b>	Indigo	6E-477	Chennai	Early_Morning	zero	Morning	Bangalore	Economy	0.83	48	1443
<b>197628</b>	Indigo	6E-6137	Chennai	Morning	zero	Morning	Bangalore	Economy	0.83	48	1443
<b>197712</b>	Indigo	6E-987	Chennai	Early_Morning	zero	Early_Morning	Bangalore	Economy	0.83	49	1443
<b>197713</b>	Indigo	6E-477	Chennai	Early_Morning	zero	Morning	Bangalore	Economy	0.83	49	1443
<b>197724</b>	Indigo	6E-6137	Chennai	Morning	zero	Morning	Bangalore	Economy	0.83	49	1549

```
In [9]: df[df['price'] == df['price'].max()]
```

```
Out[9]:
```

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
<b>261377</b>	Vistara	UK-772	Kolkata	Morning	one	Night	Delhi	Business	13.5	3	123071

```
In [10]: df[df['price'] == df['price'].min()]
```

Out[10]:

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
<b>203807</b>	AirAsia	I5-517	Chennai	Morning	zero	Morning	Hyderabad	Economy	1.17	16	1105
<b>203808</b>	GO_FIRST	G8-505	Chennai	Evening	zero	Evening	Hyderabad	Economy	1.25	16	1105
<b>203908</b>	AirAsia	I5-517	Chennai	Morning	zero	Morning	Hyderabad	Economy	1.17	17	1105
<b>203909</b>	GO_FIRST	G8-505	Chennai	Evening	zero	Evening	Hyderabad	Economy	1.25	17	1105
<b>204003</b>	AirAsia	I5-517	Chennai	Morning	zero	Morning	Hyderabad	Economy	1.17	18	1105
...	...	...	...	...	...	...	...	...	...	...	...
<b>206601</b>	Indigo	6E-7261	Chennai	Morning	one	Evening	Hyderabad	Economy	7.92	49	1105
<b>206602</b>	Indigo	6E-611	Chennai	Evening	one	Late_Night	Hyderabad	Economy	8.25	49	1105
<b>206603</b>	Indigo	6E-581	Chennai	Morning	one	Evening	Hyderabad	Economy	9.17	49	1105
<b>206604</b>	Indigo	6E-7127	Chennai	Afternoon	one	Night	Hyderabad	Economy	9.50	49	1105
<b>206605</b>	Indigo	6E-7261	Chennai	Morning	one	Night	Hyderabad	Economy	10.08	49	1105

300 rows × 11 columns

```
In [11]: df.groupby('airline')['price'].max()
```

```
Out[11]: airline
AirAsia      31917
Air_India    90970
GO_FIRST     32803
Indigo        31952
SpiceJet      34158
Vistara      123071
Name: price, dtype: int64
```

```
In [12]: df.groupby('airline')['price'].min()
```

```
Out[12]: airline
AirAsia      1105
Air_India     1526
GO_FIRST     1105
Indigo        1105
SpiceJet      1106
Vistara       1714
Name: price, dtype: int64
```

## CASE STUDY


### 1.What are the airlines in the dataset, accompanied by their frequencies?

```
In [13]: df['airline'].value_counts()
```

```
Out[13]: airline
Vistara      127859
Air_India    80892
Indigo       43120
GO_FIRST     23173
AirAsia      16098
SpiceJet      9011
Name: count, dtype: int64
```

```
In [69]: airline_counts = df['airline'].value_counts()
sns.set_style('darkgrid')
plt.figure(figsize=(25,7))
sns.barplot(x=airline_counts.index, y=airline_counts.values, palette='Blues')
plt.xlabel('AIRLINE', fontsize=16)
plt.ylabel('COUNT', fontsize=16)
plt.title('FLIGHTS BY AIRLINE', fontsize=14)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.show()
```





## 2.Show Bar Graphs representing the Departure Time & Arrival Time.

```
In [15]: df['departure_time'].value_counts()
```

```
Out[15]: departure_time
Morning      71146
Early_Morning 66790
Evening      65102
Night        48015
Afternoon    47794
Late_Night   1306
Name: count, dtype: int64
```

```
In [16]: df['arrival_time'].value_counts()
```

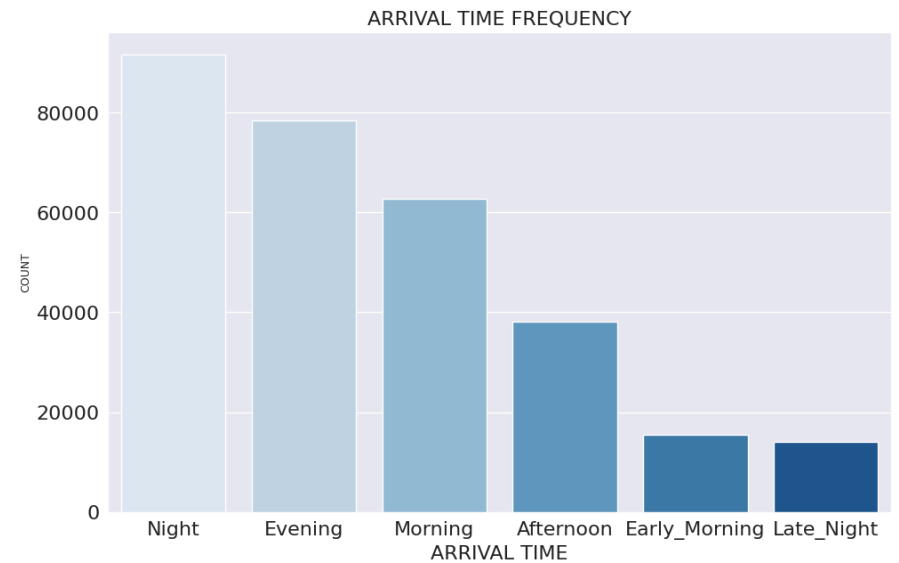

```
Out[16]: arrival_time
Night      91538
Evening    78323
Morning    62735
Afternoon  38139
Early_Morning 15417
Late_Night 14001
Name: count, dtype: int64
```

```
In [70]: departure_time = df['departure_time'].value_counts()
arrival_time = df['arrival_time'].value_counts()

sns.set_style('darkgrid')

plt.figure(figsize=(25,7))
plt.subplot(1,2,1)
sns.barplot(x=departure_time.index, y=departure_time.values, palette='Blues')
plt.xlabel('DEPARTURE TIME', fontsize=16)
plt.ylabel('COUNT', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.subplot(1,2,1)
plt.title('DEPARTURE TIME FREQUENCY', fontsize=16)

plt.subplot(1,2,2)
sns.barplot(x=arrival_time.index, y=arrival_time.values, palette='Blues')
plt.xlabel('ARRIVAL TIME', fontsize=16)
plt.ylabel('COUNT', fontsize=9)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.subplot(1,2,2)
plt.title('ARRIVAL TIME FREQUENCY', fontsize=16)
plt.show()
```



### 3.Show Bar Graphs representing the Source City & Destination City

```
In [18]: df['source_city'].value_counts()
```

```
Out[18]: source_city
Delhi      61343
Mumbai     60896
Bangalore  52061
Kolkata    46347
Hyderabad  40806
Chennai    38700
Name: count, dtype: int64
```

```
In [19]: df['destination_city'].value_counts()
```

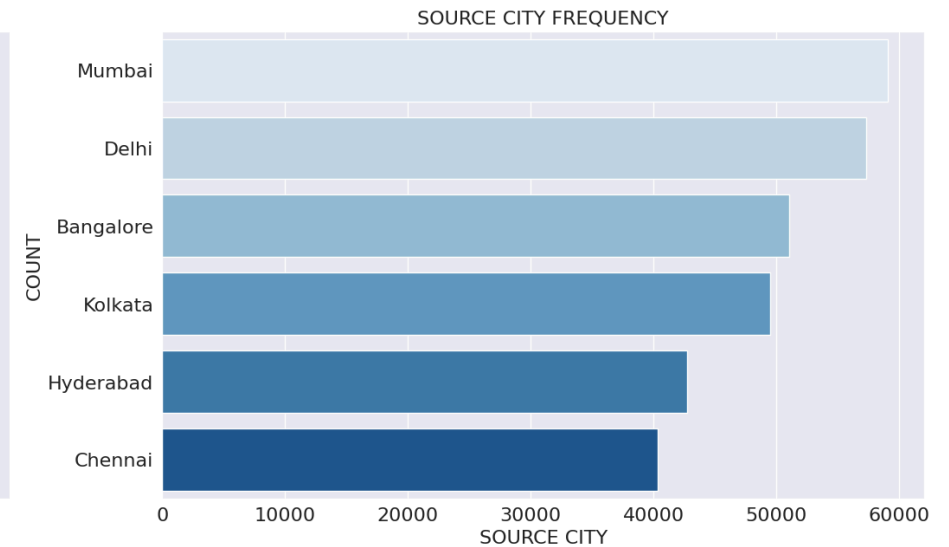
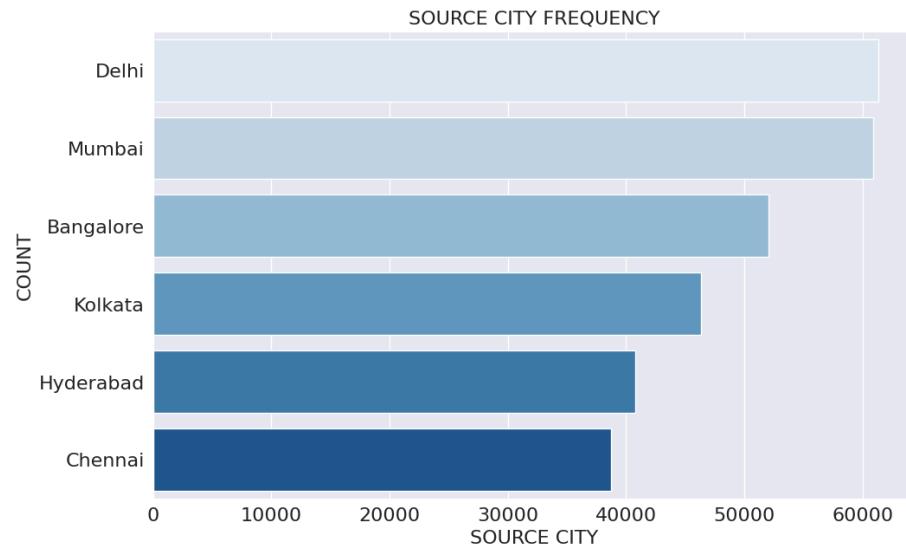
```
Out[19]: destination_city
Mumbai      59097
Delhi       57360
Bangalore   51068
Kolkata     49534
Hyderabad   42726
Chennai     40368
Name: count, dtype: int64
```

```
In [72]: source_city = df['source_city'].value_counts()
destination_city = df['destination_city'].value_counts()

plt.figure(figsize=(25,7))
plt.subplot(1,2,1)
sns.barplot(x=source_city.values, y=source_city.index, palette='Blues')
plt.xlabel('SOURCE CITY', fontsize=16)
plt.ylabel('COUNT', fontsize=16)
plt.title('SOURCE CITY FREQUENCY', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)

plt.subplot(1,2,2)
sns.barplot(x=destination_city.values, y=destination_city.index, palette='Blues')
plt.xlabel('SOURCE CITY', fontsize=16)
plt.ylabel('COUNT', fontsize=16)
plt.title('SOURCE CITY FREQUENCY', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)

plt.show()
```




## 4.Does price varies with airlines?

```
In [21]: df.groupby('airline')['price'].mean()
```

```
Out[21]: airline
AirAsia      4091.072742
Air_India    23507.019112
GO_FIRST     5652.007595
Indigo        5324.216303
SpiceJet      6179.278881
Vistara      30396.536302
Name: price, dtype: float64
```

```
In [84]: sns.catplot(x='airline', y='price', kind='bar', palette='Blues', data=df, hue='class', height=7, aspect=3)
plt.xlabel('AIRLINE', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('AVERAGE PRICE BY AIRLINE', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.show()
```



## 5.Does ticket price change based on the departure time and arrival time?

```
In [23]: df.groupby('departure_time')['price'].mean()
```

```
Out[23]: departure_time
Afternoon    18179.203331
Early_Morning  20370.676718
Evening      21232.361894
Late_Night   9295.299387
Morning      21630.760254
Night        23062.146808
Name: price, dtype: float64
```

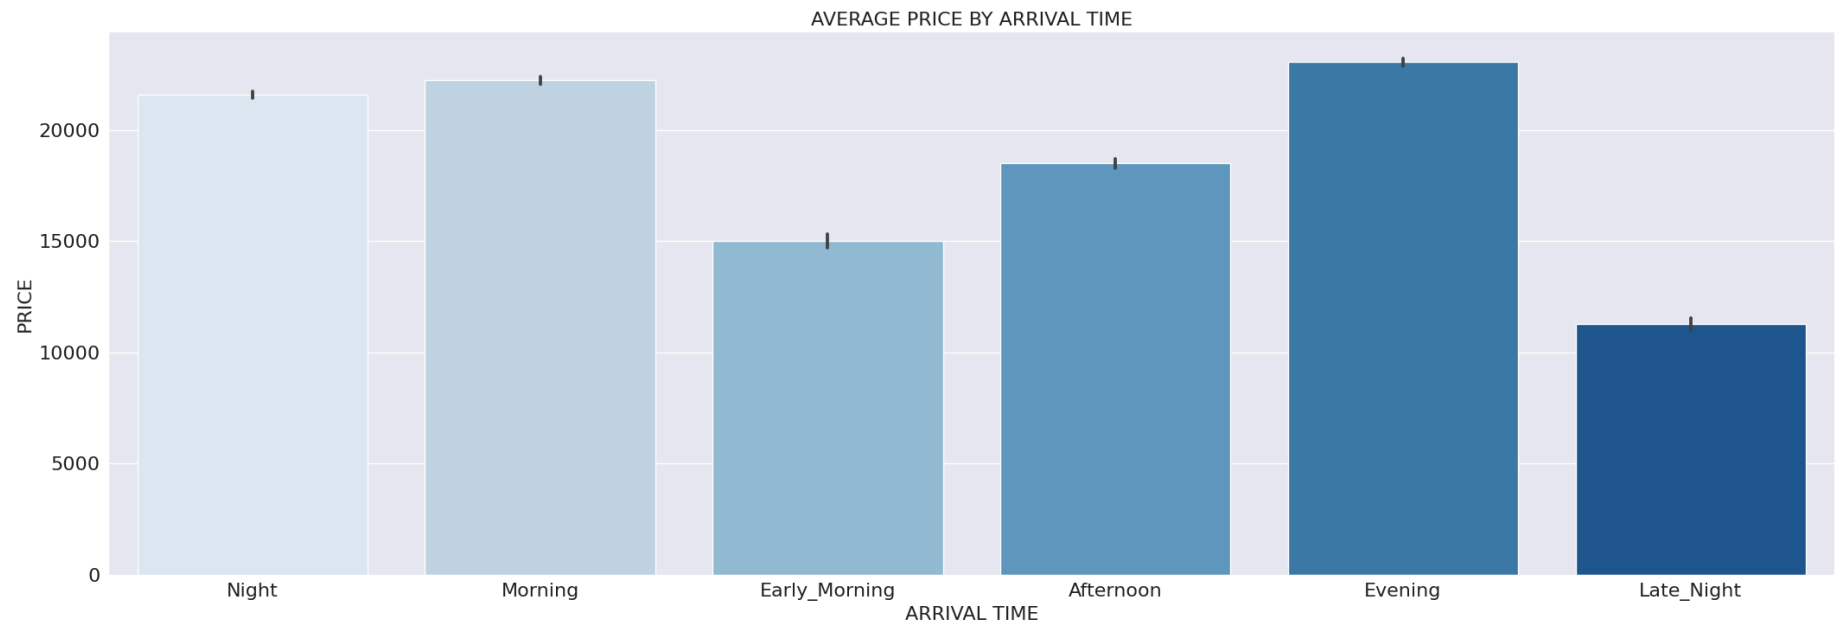
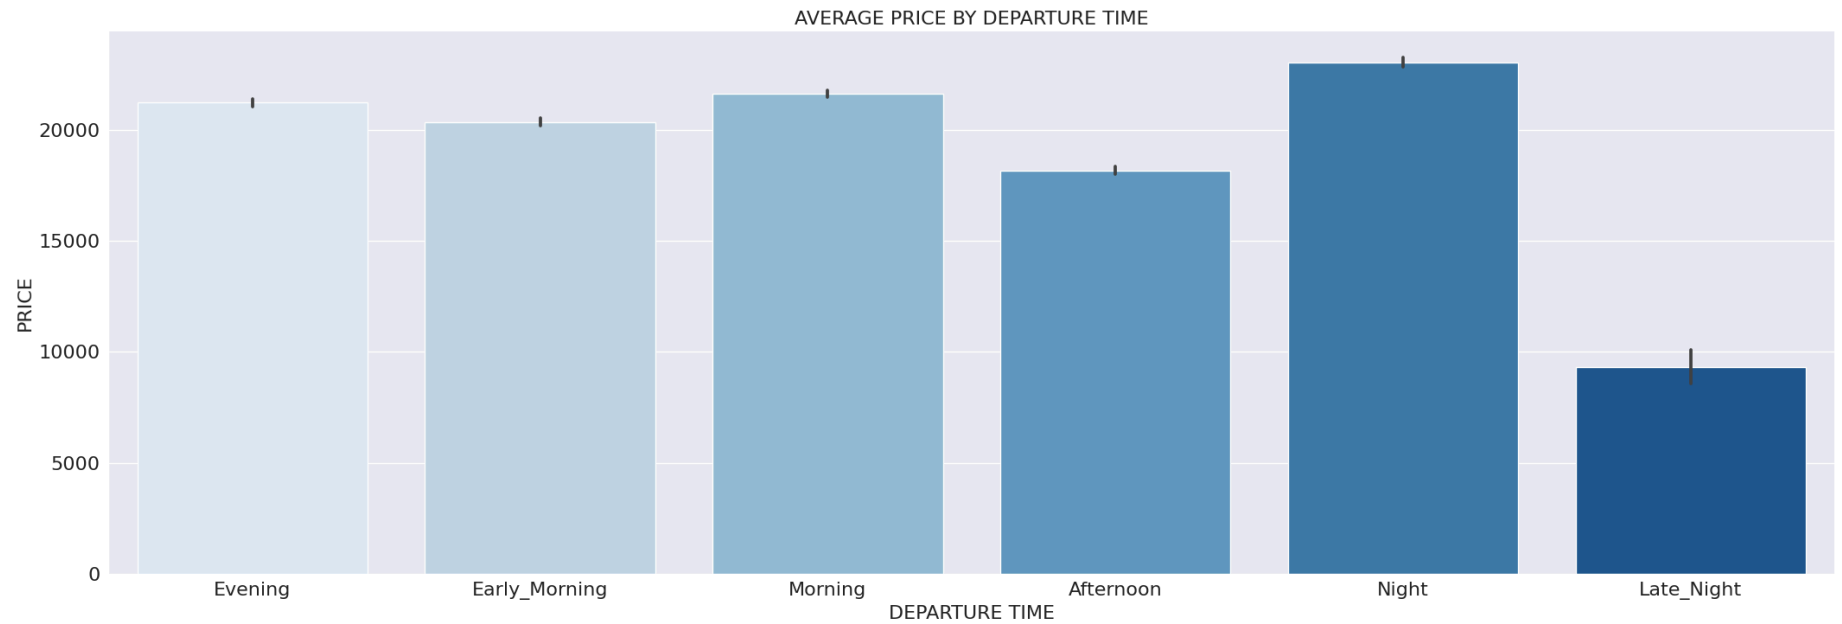
```
In [24]: df.groupby('arrival_time')['price'].mean()
```

```
Out[24]: arrival_time
Afternoon      18494.598993
Early_Morning  14993.139521
Evening        23044.371615
Late_Night     11284.906078
Morning        22231.076098
Night          21586.758341
Name: price, dtype: float64
```

```
In [75]: sns.catplot(x='departure_time', y='price', kind='bar', data=df, palette='Blues', height=7, aspect=3)
plt.xlabel('DEPARTURE TIME', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('AVERAGE PRICE BY DEPARTURE TIME', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)

sns.catplot(x='arrival_time', y='price', kind='bar', data=df, palette='Blues', height=7, aspect=3)
plt.xlabel('ARRIVAL TIME', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('AVERAGE PRICE BY ARRIVAL TIME', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.show()

plt.show()
```



**6.How the price changes with change in Source and Destination?**



```
In [26]: df.groupby('source_city')['price'].mean()
```



```
Out[26]: source_city
Bangalore    21469.460575
Chennai      21995.339871
Delhi        18951.326639
Hyderabad    20155.623879
Kolkata      21746.235679
Mumbai       21483.818839
Name: price, dtype: float64
```

```
In [27]: df.groupby('destination_city')['price'].mean()
```

```
Out[27]: destination_city
Bangalore    21593.955784
Chennai      21953.323969
Delhi        18436.767870
Hyderabad    20427.661284
Kolkata      21959.557556
Mumbai       21372.529469
Name: price, dtype: float64
```

```
In [76]: sns.catplot(x='source_city', y='price', kind='bar', data=df, palette='Blues', height=7, aspect=3)
plt.xlabel('SOURCE CITY', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('AVERAGE PRICE BY SOURCE CITY', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)

sns.catplot(x='destination_city', y='price', kind='bar', data=df, palette='Blues', height=7, aspect=3)
plt.xlabel('DESTINATION CITY', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('AVERAGE PRICE BY DESTINATION CITY', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.show()
```




**7.How is the price affected when tickets are bought in just 1 or 2 days before departure?**

```
In [29]: media_por_dia = df[df['days_left'].isin([1, 2]).groupby('days_left')['price'].mean().reset_index()
print(media_por_dia)
```

	days_left	price
0	1	21591.867151
1	2	30211.299801

```
In [77]: sns.catplot(x='days_left', y='price', kind='bar', data=media_por_dia, palette='Blues', height=7, aspect=3)
plt.xlabel('DAYS LEFT', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('AVERAGE PRICE BY 1 and 2 DAYS LEFT', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.show()
```




```
In [31]: df.groupby('days_left')['price'].mean()
```

```
Out[31]: days_left
1      21591.867151
2      30211.299801
3      28976.083569
4      25730.905653
5      26679.773368
6      24856.493902
7      25588.367351
8      24895.883995
9      25726.246072
10     25572.819134
11     22990.656070
12     22505.803322
13     22498.885384
14     22678.002363
15     21952.540852
16     20503.546237
17     20386.353949
18     19987.445168
19     19507.677375
20     19699.983390
21     19430.494058
22     19590.667385
23     19840.913451
24     19803.908896
25     19571.641791
26     19238.290278
27     19950.866195
28     19534.986047
29     19744.653119
30     19567.580834
31     19392.706612
32     19258.135308
33     19306.271739
34     19562.008266
35     19255.652996
36     19517.688444
37     19506.306516
38     19734.912316
39     19262.095556
```

```
40    19144.972439
41    19347.440460
42    19154.261659
43    19340.528894
44    19049.080174
45    19199.876307
46    19305.351623
47    18553.272038
48    18998.126851
49    18992.971888
Name: price, dtype: float64
```

```
In [78]: plt.figure(figsize=(25,7))
sns.barplot(x='days_left', y='price', data=df, palette='Blues')
plt.xlabel('DAYS LEFT', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('AVERAGE PRICE BY DAYS LEFT', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.show()
```




**8.How does the ticket price vary between Economy and Business class?**

```
In [33]: df.groupby('class')['price'].mean()
```

```
Out[33]: class
Business    52540.081124
Economy      6572.342383
Name: price, dtype: float64
```

```
In [80]: plt.figure(figsize=(25,7))
sns.barplot(x='class', y='price', data=df, palette='Blues', hue='class')
plt.xlabel('CLASS', fontsize=16)
plt.ylabel('PRICE', fontsize=16)
plt.title('PRICE BY CLASS', fontsize=16)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.legend(fontsize=12)
plt.show()
```



**9.What will be the Average Price of Vistara airline for a flight from Delhi to Hyderabad in Business Class?**

```
In [35]: question9_df = df [(df['airline'] == 'Vistara') & (df['source_city'] == 'Delhi') & (df['destination_city'] == 'Hyderabad') & (
question9_df
```

```
Out[35]:
```

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
<b>219123</b>	Vistara	UK-871	Delhi	Night	zero	Night	Hyderabad	Business	2.33	1	30630
<b>219124</b>	Vistara	UK-879	Delhi	Evening	zero	Evening	Hyderabad	Business	2.25	1	38470
<b>219129</b>	Vistara	UK-955	Delhi	Evening	one	Night	Hyderabad	Business	27.17	1	63513
<b>219130</b>	Vistara	UK-955	Delhi	Evening	one	Afternoon	Hyderabad	Business	18.50	1	65764
<b>219131</b>	Vistara	UK-985	Delhi	Evening	one	Night	Hyderabad	Business	25.08	1	69113
...	...	...	...	...	...	...	...	...	...	...	...
<b>221863</b>	Vistara	UK-963	Delhi	Morning	one	Early_Morning	Hyderabad	Business	23.00	49	53937
<b>221864</b>	Vistara	UK-985	Delhi	Evening	one	Early_Morning	Hyderabad	Business	12.00	49	59537
<b>221865</b>	Vistara	UK-985	Delhi	Evening	one	Afternoon	Hyderabad	Business	16.42	49	59537
<b>221866</b>	Vistara	UK-955	Delhi	Evening	one	Early_Morning	Hyderabad	Business	14.08	49	61889
<b>221867</b>	Vistara	UK-955	Delhi	Evening	one	Afternoon	Hyderabad	Business	18.50	49	61889

1660 rows × 11 columns

```
In [36]: price = round(question9_df['price'].mean(), 0)
print ('The Average Price of Vistara airline for a flight from Delhi to Hyderabad in Business Class is:', price)
```

The Average Price of Vistara airline for a flight from Delhi to Hyderabad in Business Class is: 47940.0

## FEATURES ENGINEERING

```
In [37]: df.head()
```

```
Out[37]:
```

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_left	price
0	SpiceJet	SG-8709	Delhi	Evening	zero	Night	Mumbai	Economy	2.17	1	5953
1	SpiceJet	SG-8157	Delhi	Early_Morning	zero	Morning	Mumbai	Economy	2.33	1	5953
2	AirAsia	I5-764	Delhi	Early_Morning	zero	Early_Morning	Mumbai	Economy	2.17	1	5956
3	Vistara	UK-995	Delhi	Morning	zero	Afternoon	Mumbai	Economy	2.25	1	5955
4	Vistara	UK-963	Delhi	Morning	zero	Morning	Mumbai	Economy	2.33	1	5955

```
In [38]: df_new = df.copy()
```

```
In [39]: df_new.drop(columns= 'flight', inplace=True)
```

```
In [40]: df_new['stops'].value_counts()
```


```
Out[40]: stops
one          250863
zero         36004
two_or_more  13286
Name: count, dtype: int64
```

```
In [41]: df_new['departure_time'].value_counts()
```

```
Out[41]: departure_time
Morning          71146
Early_Morning    66790
Evening          65102
Night            48015
Afternoon        47794
Late_Night       1306
Name: count, dtype: int64
```

```
In [42]: corr = df_new.select_dtypes(include=[np.number]).corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap of Numeric Variables')
plt.show()
```





```
In [43]: # Diccionario de mapeo
franjas = {
    "Late_Night": "02:00:00",
    "Early_Morning": "05:00:00",
    "Morning": "09:00:00",
    "Afternoon": "15:00:00",
    "Evening": "19:00:00",
    "Night": "22:00:00"
}

# Reemplazar con la hora correspondiente
df_new["departure_time"] = df_new["departure_time"].map(franjas)
```

```
# Convertir a tipo time
df_new["departure_time"] = pd.to_datetime(df_new["departure_time"], format="%H:%M:%S").dt.time
```

```
In [44]: # Reemplazar con la hora correspondiente
df_new["arrival_time"] = df_new["arrival_time"].map(franjas)

# Convertir a tipo time
df_new["arrival_time"] = pd.to_datetime(df_new["arrival_time"], format="%H:%M:%S").dt.time
```

```
In [45]: # Encoding cíclico para las franjas horarias
def encode_time(df_new, col):
    seconds = df_new[col].apply(lambda t: t.hour*3600 + t.minute*60 + t.second)
    df_new[col+"_sin"] = np.sin(2 * np.pi * seconds / 86400)
    df_new[col+"_cos"] = np.cos(2 * np.pi * seconds / 86400)
    return df_new

df_new = encode_time(df_new, "departure_time")
df_new = encode_time(df_new, "arrival_time")
```

```
In [46]: # Definimos las variables con más y menos clases para el codificador
low_card = ["airline", "stops", "class"] # pocas clases → OHE
high_card = ["source_city", "destination_city"] # muchas clases → Target Encoding

preprocessor = ColumnTransformer(
    transformers=[
        ("low_card_ohe", OneHotEncoder(handle_unknown="ignore"), low_card),
        ("high_card_target", TargetEncoder(), high_card)
    ],
    remainder="passthrough" # deja el resto de columnas tal cual
)

pipeline = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("model", RandomForestRegressor()) # pon aquí tu modelo
])
```

```
In [47]: df_new.drop(columns='departure_time', inplace=True)
```


```
In [48]: df_new.drop(columns='arrival_time', inplace=True)
```

```
In [49]: # numéricas/passthrough (ajusta si tienes más)
num_cols = [c for c in df_new.columns
             if c not in low_card + high_card + ["price"]] # quita la y si está en X
```

```
In [50]: # 2) Crea los transformadores
ohe = OneHotEncoder(handle_unknown="ignore", sparse_output=False) # nombres expandibles
tgt = TargetEncoder()

preprocessor = ColumnTransformer(
    transformers=[
        ("ohe_low", ohe, low_card),
        ("tgt_high", tgt, high_card),
        ("num", "passthrough", num_cols)
    ],
    verbose_feature_names_out=False
)
```

```
In [51]: # 3) Haz que devuelva un **DataFrame** (no una matriz)
preprocessor.set_output(transform="pandas")
```



```
In [52]: # 4) Separa X, y y transforma
y = df_new["price"]
X = df_new.drop(columns=["price"])

X_enc = preprocessor.fit_transform(X, y) # ← TargetEncoder necesita y
```

```
In [53]: # 5) Revisa resultado
print(X_enc.columns.tolist()[:25]) # ver primeros nombres
X_enc.head()
```

```
['airline_AirAsia', 'airline_Air_India', 'airline_GO_FIRST', 'airline_Indigo', 'airline_SpiceJet', 'airline_Vistara', 'stops_one', 'stops_two_or_more', 'stops_zero', 'class_Business', 'class_Economy', 'source_city', 'destination_city', 'duration', 'days_left', 'departure_time_sin', 'departure_time_cos', 'arrival_time_sin', 'arrival_time_cos']
```

Out[53]:

	airline_AirAsia	airline_Air_India	airline_GO_FIRST	airline_Indigo	airline_SpiceJet	airline_Vistara	stops_one	stops_two_or_more	stops_zero
0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0

## ML MODELS AND EVALUATION

```
In [54]: # Train test split
X_train, X_test, y_train, y_test = train_test_split(X_enc, y, test_size=0.2, random_state=42)

# Entrenamos RandomForestRegressor
model = RandomForestRegressor(random_state=42)
model.fit(X_train, y_train)

# Predicción
preds = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, preds))
# MAE (Mean Absolute Error)
mae = mean_absolute_error(y_test, preds)

# R² (Coeficiente de determinación)
r2 = r2_score(y_test, preds)

print(f"RMSE: {rmse:,.2f}")
print(f"MAE: {mae:,.2f}")
print(f"R²: {r2:,.4f}")
```

RMSE: 2,774.16  
MAE: 1,075.83  
R<sup>2</sup>: 0.9851

```
In [55]: # Modelo base
xgb = XGBRegressor(random_state=42, n_jobs=-1)

# Espacio de búsqueda
param_dist = {
    "n_estimators": [200, 400, 600, 800],
    "learning_rate": [0.01, 0.05, 0.1, 0.2],
    "max_depth": [3, 5, 6, 8, 10],
    "subsample": [0.6, 0.8, 1.0],
    "colsample_bytree": [0.6, 0.8, 1.0],
    "gamma": [0, 1, 5]
}

# Búsqueda aleatoria
random_search = RandomizedSearchCV(
    estimator=xgb,
    param_distributions=param_dist,
    n_iter=20,          # número de combinaciones a probar
    scoring="neg_mean_squared_error",
    cv=3,
    verbose=1,
    random_state=42,
    n_jobs=-1
)

random_search.fit(X_train, y_train)

# Mejor modelo
best_xgb = random_search.best_estimator_
print("Mejores parámetros:", random_search.best_params_)

# Predicciones
preds = best_xgb.predict(X_test)

# Métricas
rmse = np.sqrt(mean_squared_error(y_test, preds))
mae = mean_absolute_error(y_test, preds)
```

```

r2 = r2_score(y_test, preds)

print(f"RMSE: {rmse:,.2f}")
print(f"MAE: {mae:,.2f}")
print(f"R²: {r2:,.4f}")

```

Fitting 3 folds for each of 20 candidates, totalling 60 fits

Mejores parámetros: {'subsample': 1.0, 'n\_estimators': 600, 'max\_depth': 8, 'learning\_rate': 0.2, 'gamma': 1, 'colsample\_bytree': 0.8}

RMSE: 2,458.58

MAE: 1,258.63

R²: 0.9883

```

In [56]: # XGBoost Cross Validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

rmse_scores = []
mae_scores = []
r2_scores = []

for train_idx, test_idx in kf.split(X_enc):
    X_train, X_test = X_enc.iloc[train_idx], X_enc.iloc[test_idx]
    y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]

    # Entrenar
    best_xgb.fit(X_train, y_train)

    # Predecir
    preds = best_xgb.predict(X_test)

    # Métricas
    rmse_scores.append(np.sqrt(mean_squared_error(y_test, preds)))
    mae_scores.append(mean_absolute_error(y_test, preds))
    r2_scores.append(r2_score(y_test, preds))

# Resultados
print("📊 XGBoost")
print("RMSE por fold:", rmse_scores)
print("MAE por fold:", mae_scores)
print("R² por fold:", r2_scores)

```

```

print("\nPromedios:")
print(f"RMSE medio: {np.mean(rmse_scores):.2f} ± {np.std(rmse_scores):.2f}")
print(f"MAE medio:  {np.mean(mae_scores):.2f} ± {np.std(mae_scores):.2f}")
print(f"R² medio:   {np.mean(r2_scores):.4f} ± {np.std(r2_scores):.4f}")

```



XGBoost

RMSE por fold: [2458.581620710177, 2430.2753569624438, 2455.7461470454214, 2473.094920448661, 2429.885715398126]  
 MAE por fold: [1258.631701793551, 1251.4457718248566, 1274.2230557450307, 1257.2291140699792, 1257.9916161314281]  
 R² por fold: [0.9882738478205182, 0.9885573174488035, 0.9882888347373722, 0.9881418013929509, 0.9885031376654303]

Promedios:

RMSE medio: 2,449.52 ± 16.93

MAE medio: 1,259.90 ± 7.60

R² medio: 0.9884 ± 0.0002

```

In [57]: # RandomForestRegressor Cross Validation
rf_model = RandomForestRegressor(
    n_estimators=500,
    max_depth=None,
    random_state=42,
    n_jobs=-1
)

# Cross Validation (5 folds)
kf = KFold(n_splits=5, shuffle=True, random_state=42)

rmse_scores = []
mae_scores = []
r2_scores = []

for train_idx, test_idx in kf.split(X_enc):
    X_train, X_test = X_enc.iloc[train_idx], X_enc.iloc[test_idx]
    y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]

    rf_model.fit(X_train, y_train)
    preds = rf_model.predict(X_test)

    rmse_scores.append(np.sqrt(mean_squared_error(y_test, preds)))
    mae_scores.append(mean_absolute_error(y_test, preds))
    r2_scores.append(r2_score(y_test, preds))

```

```

print("📊 RandomForest")
print("RMSE por fold:", rmse_scores)
print("MAE por fold:", mae_scores)
print("R² por fold:", r2_scores)

print("\nPromedios:")
print(f"RMSE medio: {np.mean(rmse_scores):.2f} ± {np.std(rmse_scores):.2f}")
print(f"MAE medio: {np.mean(mae_scores):.2f} ± {np.std(mae_scores):.2f}")
print(f"R² medio: {np.mean(r2_scores):.4f} ± {np.std(r2_scores):.4f}")

```

📊 RandomForest

RMSE por fold: [2764.2436860423827, 2730.88367323717, 2710.4547542572122, 2717.4439348411806, 2679.97774849615]  
MAE por fold: [1071.0650716570694, 1061.0252021783094, 1076.4131363122488, 1060.455053690341, 1054.8170366800414]  
R² por fold: [0.9851769041388689, 0.9855514825974288, 0.9857334983524717, 0.9856827920038108, 0.986014757172144]

Promedios:

RMSE medio: 2,720.60 ± 27.47

MAE medio: 1,064.76 ± 7.83

R² medio: 0.9856 ± 0.0003

In [106...

```

# MAE vs Precio Real
rf_model.fit(X_train, y_train)
rf_preds = rf_model.predict(X_test)

best_xgb.fit(X_train, y_train)
xgb_preds = best_xgb.predict(X_test)

# Errores absolutos
rf_errors = np.abs(rf_preds - y_test)
xgb_errors = np.abs(xgb_preds - y_test)

# DataFrame para seaborn
import pandas as pd
df_plot = pd.DataFrame({
    "Precio_real": y_test,
    "Error": rf_errors,
    "Modelo": "RandomForest"
})
df_plot = pd.concat([
    df_plot,
    pd.DataFrame({

```




```

        "Precio_real": y_test,
        "Error": xgb_errors,
        "Modelo": "XGBoost"
    })
])

# --- Gráfico ---
plt.figure(figsize=(25,6))
sns.scatterplot(data=df_plot, x="Precio_real", y="Error", hue="Modelo", alpha=0.4)
sns.regplot(data=df_plot[df_plot["Modelo"]=="RandomForest"],
            x="Precio_real", y="Error", scatter=False, color="blue", label="Tendencia RF")
sns.regplot(data=df_plot[df_plot["Modelo"]=="XGBoost"],
            x="Precio_real", y="Error", scatter=False, color="red", label="Tendencia XGB")

plt.title("Errores absolutos vs Precio real con líneas de tendencia", fontsize=20)
plt.xlabel("Precio real", fontsize=20)
plt.ylabel("Error", fontsize=20)
plt.xticks(fontsize=20)
plt.yticks(fontsize=20)
plt.legend(fontsize=16)
plt.show()

```



```

In [105... # Datos de resultados promedio
resultados = {
    "RandomForest": {

```

```

        "RMSE": 2720.60,
        "MAE": 1064.76,
        "R2": 0.9856
    },
    "XGBoost": {
        "RMSE": 2454.20,
        "MAE": 1258.45,
        "R2": 0.9883
    }
}

# Configuración
metricas = ["RMSE", "MAE", "R2"]
x = np.arange(len(metricas)) # posiciones
ancho = 0.35 # ancho de las barras

fig, ax = plt.subplots(figsize=(25, 8))

# Barras para RF y XGB
rf_vals = [resultados["RandomForest"][m] for m in metricas]
xgb_vals = [resultados["XGBoost"][m] for m in metricas]

ax.bar(x - ancho/2, rf_vals, ancho, label="RandomForest", color="skyblue")
ax.bar(x + ancho/2, xgb_vals, ancho, label="XGBoost", color="steelblue")

# Etiquetas y formato
ax.set_ylabel("Valor", fontsize=20)
ax.set_title("Comparación de métricas entre RandomForest y XGBoost", fontsize=20)

# Ticks del eje X
ax.set_xticks(x)
ax.set_xticklabels(metricas, fontsize=20)
ax.tick_params(axis='y', labelsize=20)

ax.legend(fontsize=20)
ax.grid(axis='y', linestyle='--', alpha=0.6)

# Mostrar valores sobre las barras
for i, v in enumerate(rf_vals):
    ax.text(i - ancho/2, v + max(0.01*v, 0.005), f"{v:.2f}", ha='center', va='bottom', fontsize=20)


```

```

for i, v in enumerate(xgb_vals):
    ax.text(i + ancho/2, v + max(0.01*v, 0.005), f"{v:.2f}", ha='center', va='bottom', fontsize=20)

plt.tight_layout()
plt.show()

```



## FEATURE IMPORTANCES

```

In [107... # Feature importances series ordenadas
rf_imp = pd.Series(rf_model.feature_importances_, index=X_enc.columns).sort_values(ascending=False)
xgb_imp = pd.Series(best_xgb.feature_importances_, index=X_enc.columns).sort_values(ascending=False)

top_k = 15
fig, axs = plt.subplots(1, 2, figsize=(25, 10), sharey=True)

rf_imp.head(top_k).iloc[::-1].plot(kind="barh", ax=axs[0], title="RF - Top 15", fontsize=20, color='steelblue')
xgb_imp.head(top_k).iloc[::-1].plot(kind="barh", ax=axs[1], title="XGBoost - Top 15", fontsize=20, color='steelblue')

plt.xticks(fontsize=20)
plt.yticks(fontsize=20)
plt.tight_layout()


```

```

axs[0].set_title("RF - Top 15", fontsize=20)
axs[1].set_title("XGBoost - Top 15", fontsize=20)

plt.show()

```



```

In [108... # Feature importances (variables iniciales)
def group_importances(imp_series):
    grp = {}
    for col, val in imp_series.items():
        key = col.split("_")[0] # prefijo antes del primer "_"
        grp[key] = grp.get(key, 0.0) + val
    return pd.Series(grp).sort_values(ascending=False)

rf_grp = group_importances(rf_imp)
xgb_grp = group_importances(xgb_imp)

fig, axs = plt.subplots(1, 2, figsize=(25, 10), sharey=True)

```

```


rf_grp.head(10).iloc[::-1].plot(kind="barh", ax=axes[0], title="RF (agrupado por variable)", fontsize=20, color='steelblue')
xgb_grp.head(10).iloc[::-1].plot(kind="barh", ax=axes[1], title="XGB (agrupado por variable)", fontsize=20, color='steelblue')

plt.xticks(fontsize=20)
plt.yticks(fontsize=20)
plt.tight_layout()

axes[0].set_title("RF (agrupado por variable)", fontsize=20)
axes[1].set_title("XGB (agrupado por variable)", fontsize=20)

plt.tight_layout()
plt.show()

```



# SHAP

```
In [62]: X_shap = X_enc.sample(1000, random_state=42) # sube a 2000 si puedes
```

```
explainer = shap.TreeExplainer(best_xgb) # sin approximate  
shap_vals = explainer.shap_values(X_shap, check_additivity=True)
```

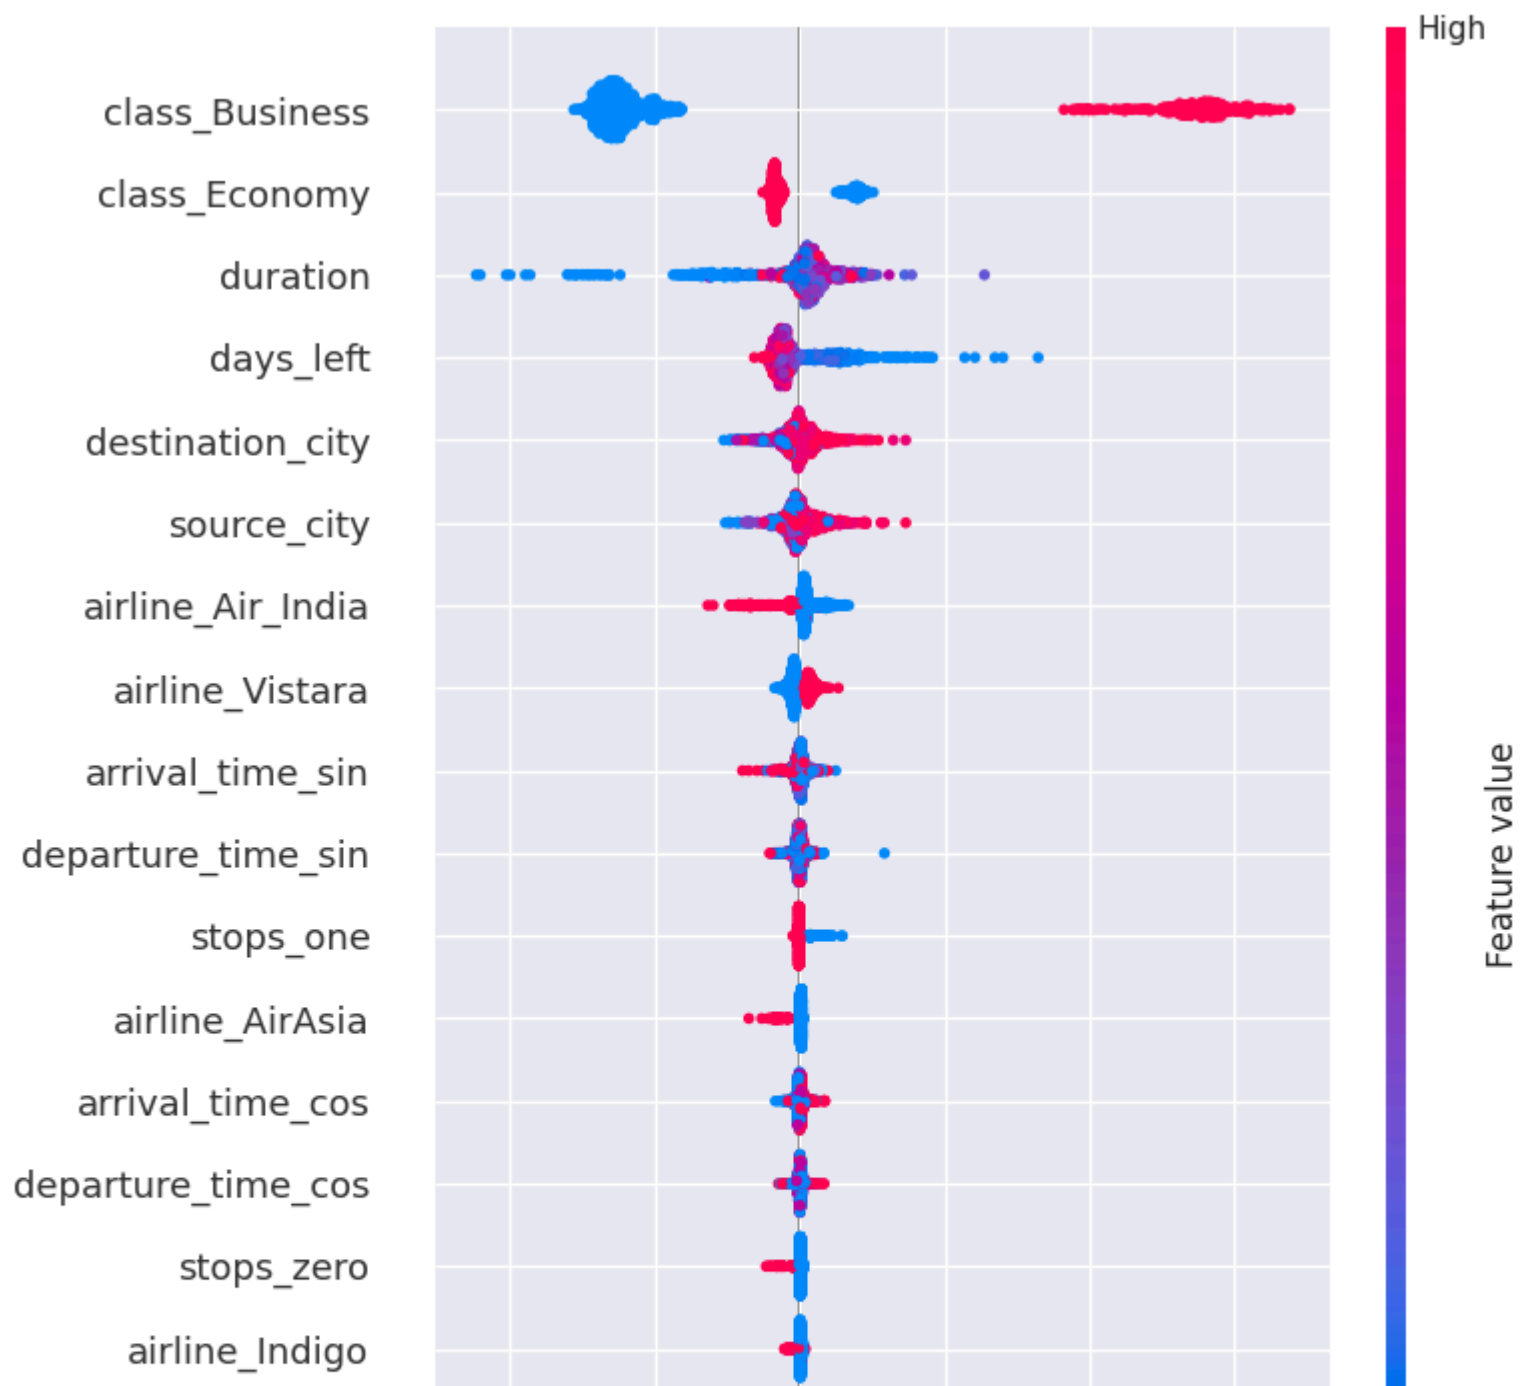
```
# Global
```


```
shap.summary_plot(shap_vals, X_shap)
```

```
# Local
```


```
i = 10
```

```
shap.waterfall_plot(  
    shap.Explanation(values=shap_vals[i],  
                    base_values=explainer.expected_value,  
                    data=X_shap.iloc[i,:],  
                    feature_names=X_shap.columns.tolist())  
)
```









## INSIGHTS

### 1. Distribution and Trends

- The dataset shows a wide variability in prices, with extreme cases on both ends (very cheap and very expensive flights).
  - The `class` variable has a clear impact: **Business** flights have significantly higher prices than **Economy**.
  - `days_left` (days remaining until the flight) follows the expected trend: the fewer the days left, the higher the price, with sharp increases in the last days before departure.
  - **Flight duration** shows an inverse relationship with price for short and direct routes, but for longer routes the relationship is not always linear, likely due to combinations of stops and cities.
- 

## 2. Model Results

### RandomForest:

- RMSE  $\approx$  **2,720**
- MAE  $\approx$  **1,065**
- $R^2 \approx$  **0.9856**
- Strong overall performance, low errors, and high explanatory power, though it tends to make bigger mistakes for high-priced flights.

### XGBoost:

- RMSE  $\approx$  **2,454**
  - MAE  $\approx$  **1,258**
  - $R^2 \approx$  **0.9883**
  - Slight improvement in RMSE and  $R^2$  compared to RF, with a small increase in MAE.
  - Cross-validation confirms model stability and generalization.
- 

## 3. Feature Importance

In both models, the most influential variables globally are:

1. **class**
2. **days\_left**

3. **duration**
4. **stops**
5. Time-related features ( `departure_time_sin` , `arrival_time_cos` )
6. Origin/destination cities (moderate importance)

Airline-related features have lower but still relevant impact.

---

## 4. Explainability

Using **SHAP**:

- Not being Business class strongly reduces the predicted price.
- Booking with very few days left significantly increases the price.
- Shorter flight duration tends to lower the price.
- Airline and departure time effects are secondary but still present.

Both models tend to have higher prediction errors for expensive flights, suggesting that extreme values in the dataset are harder to model.

---

## 5. Conclusions and Opportunities

- Both models are robust, but **XGBoost provides a slight edge** in overall accuracy.
- Performance for high-priced flights could be improved by:
  - Balancing the dataset in those price ranges.
  - Using hybrid models or transforming the target variable (e.g., `log(price)` ).
- SHAP interpretability adds value in understanding the drivers of price and could be leveraged for dynamic pricing or user recommendations.

**PLEASE UP VOTE AND SUPPORT MY WORK AS STUDENT!! :) Thanks for reading!**

<https://www.linkedin.com/in/adrianzambranaacquaroni/>

ADRIÁN ZAMBRANA | Aug 2025